

CSC 374/407: Computer Systems II

Lecture 4

Joseph Phillips
De Paul University

2017 January 31

Copyright © 2011-2017 Joseph Phillips
All rights reserved


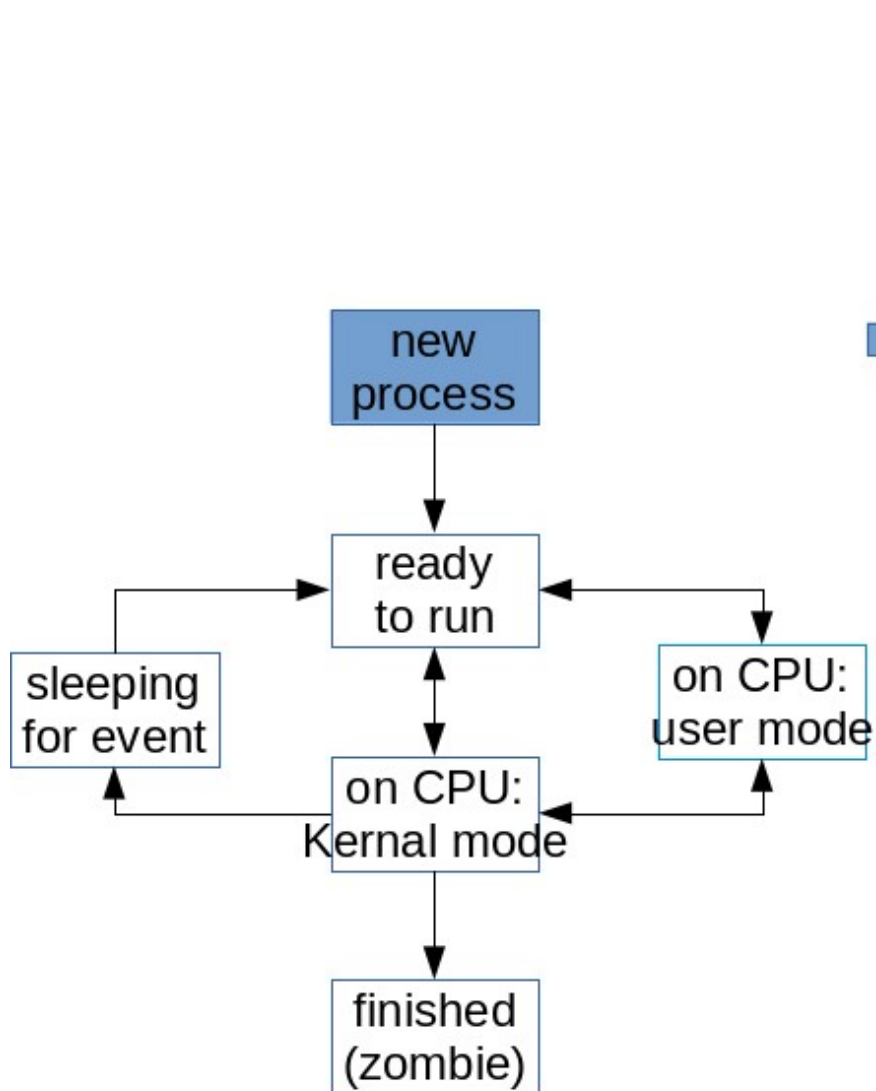
Reading

- ♦ Bryant & O'Hallaron “*Computer Systems, 3rd Ed.*”
 - ♦ Chapter 8: Exception Control Flow (8.5 Signals)
- ♦ Hoover “*System Programming*”
 - ♦ System Calls 7.1-7.4

Topics

- ◆ Signals
 - **sigaction()**
 - **kill()**
 - **alarm()**

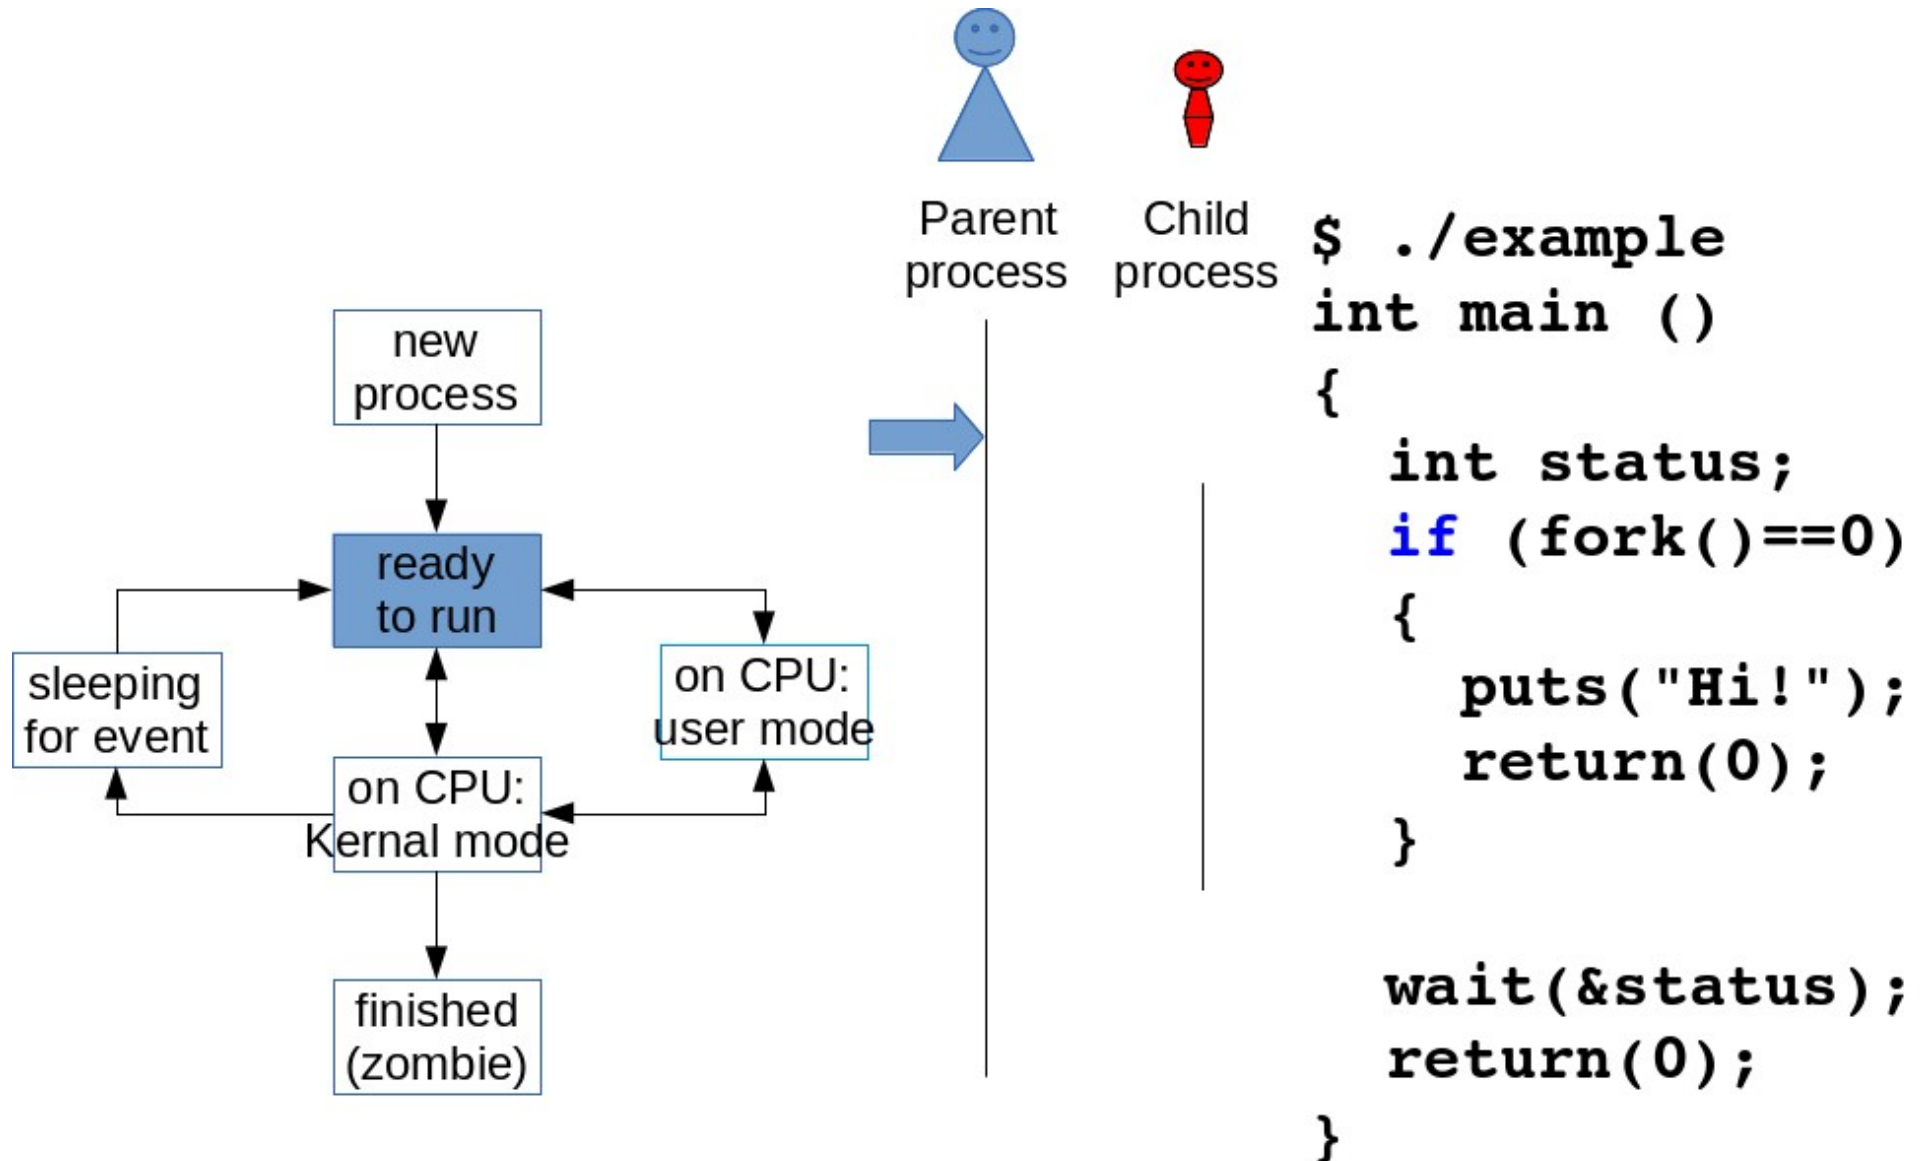
Recall fork system call (1)



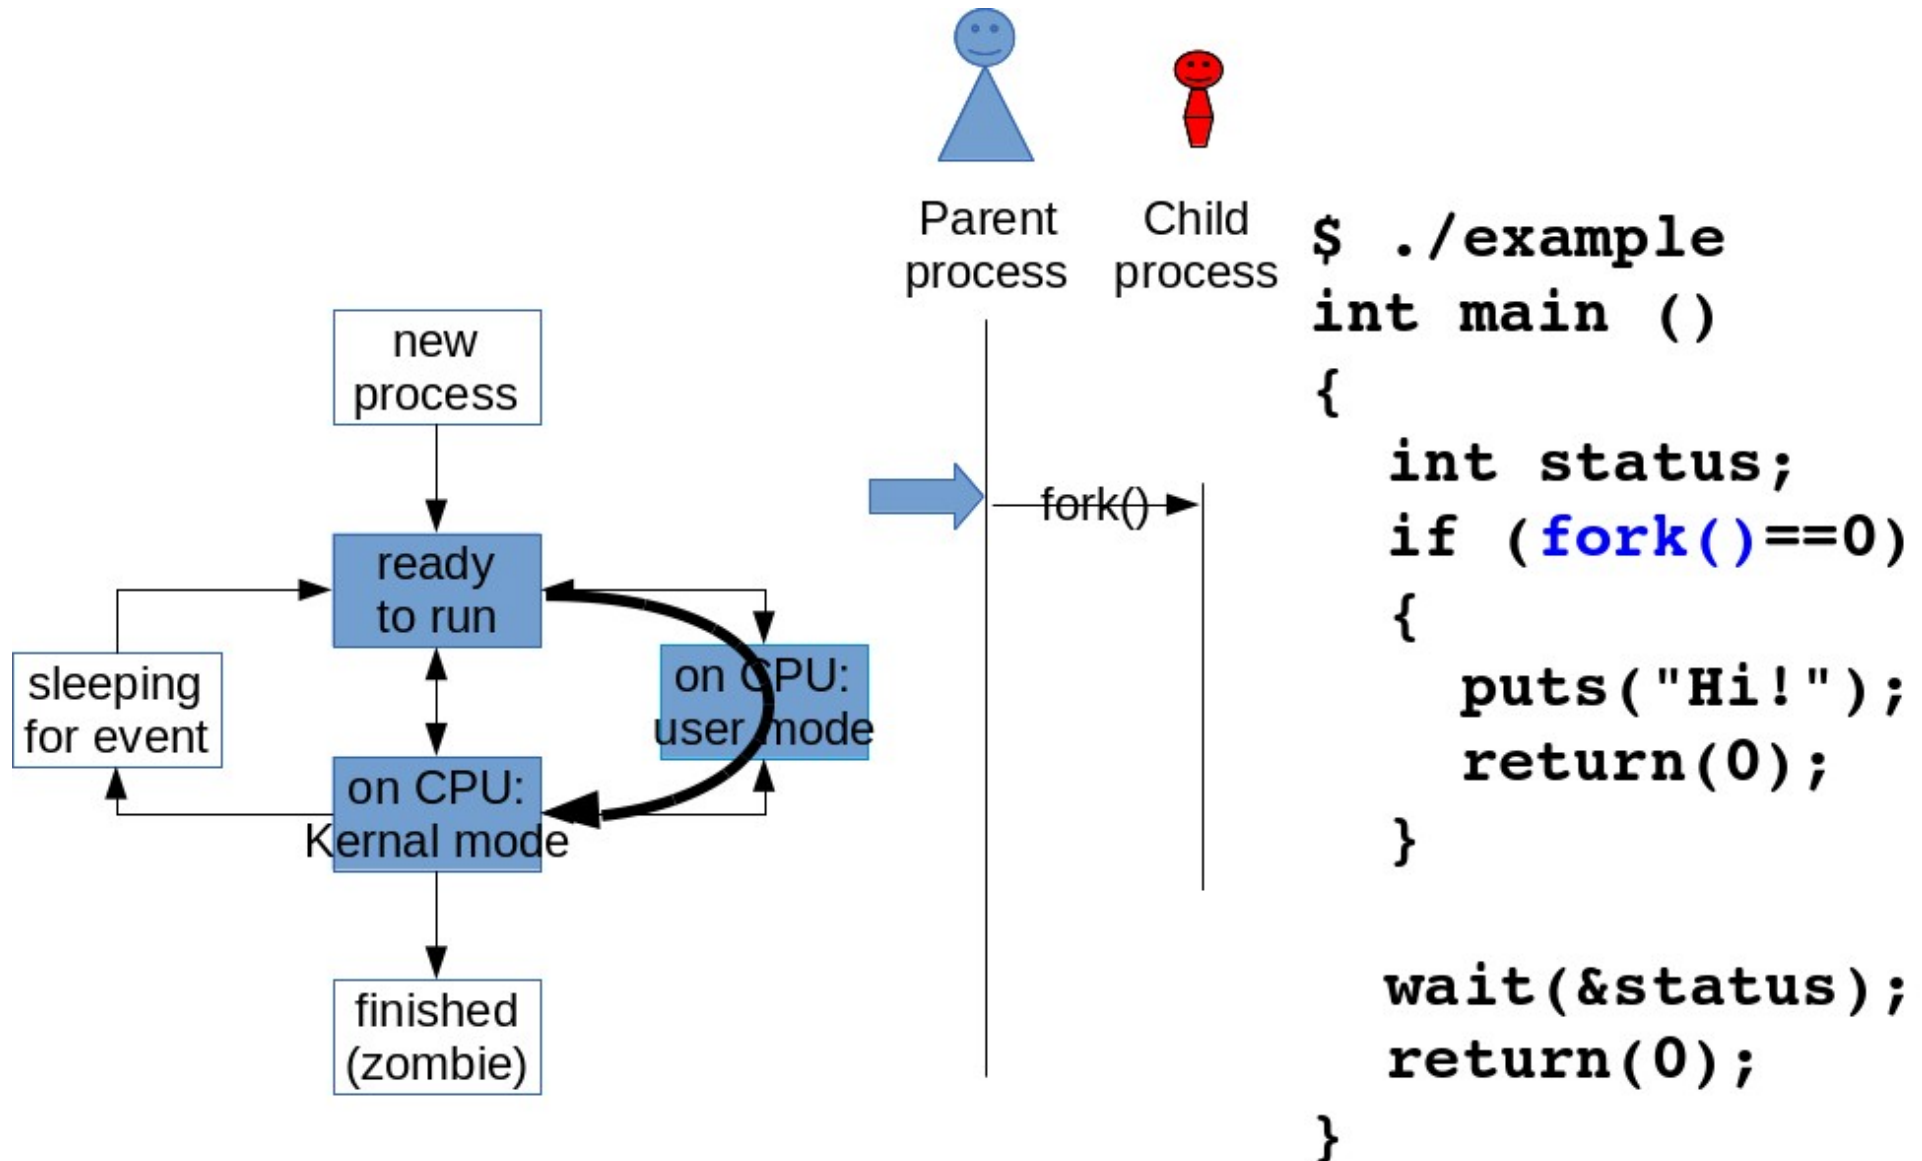
```
$ ./example
int main ()
{
    int status;
    if (fork()==0)
    {
        puts("Hi!");
        return(0);
    }

    wait(&status);
    return(0);
}
```

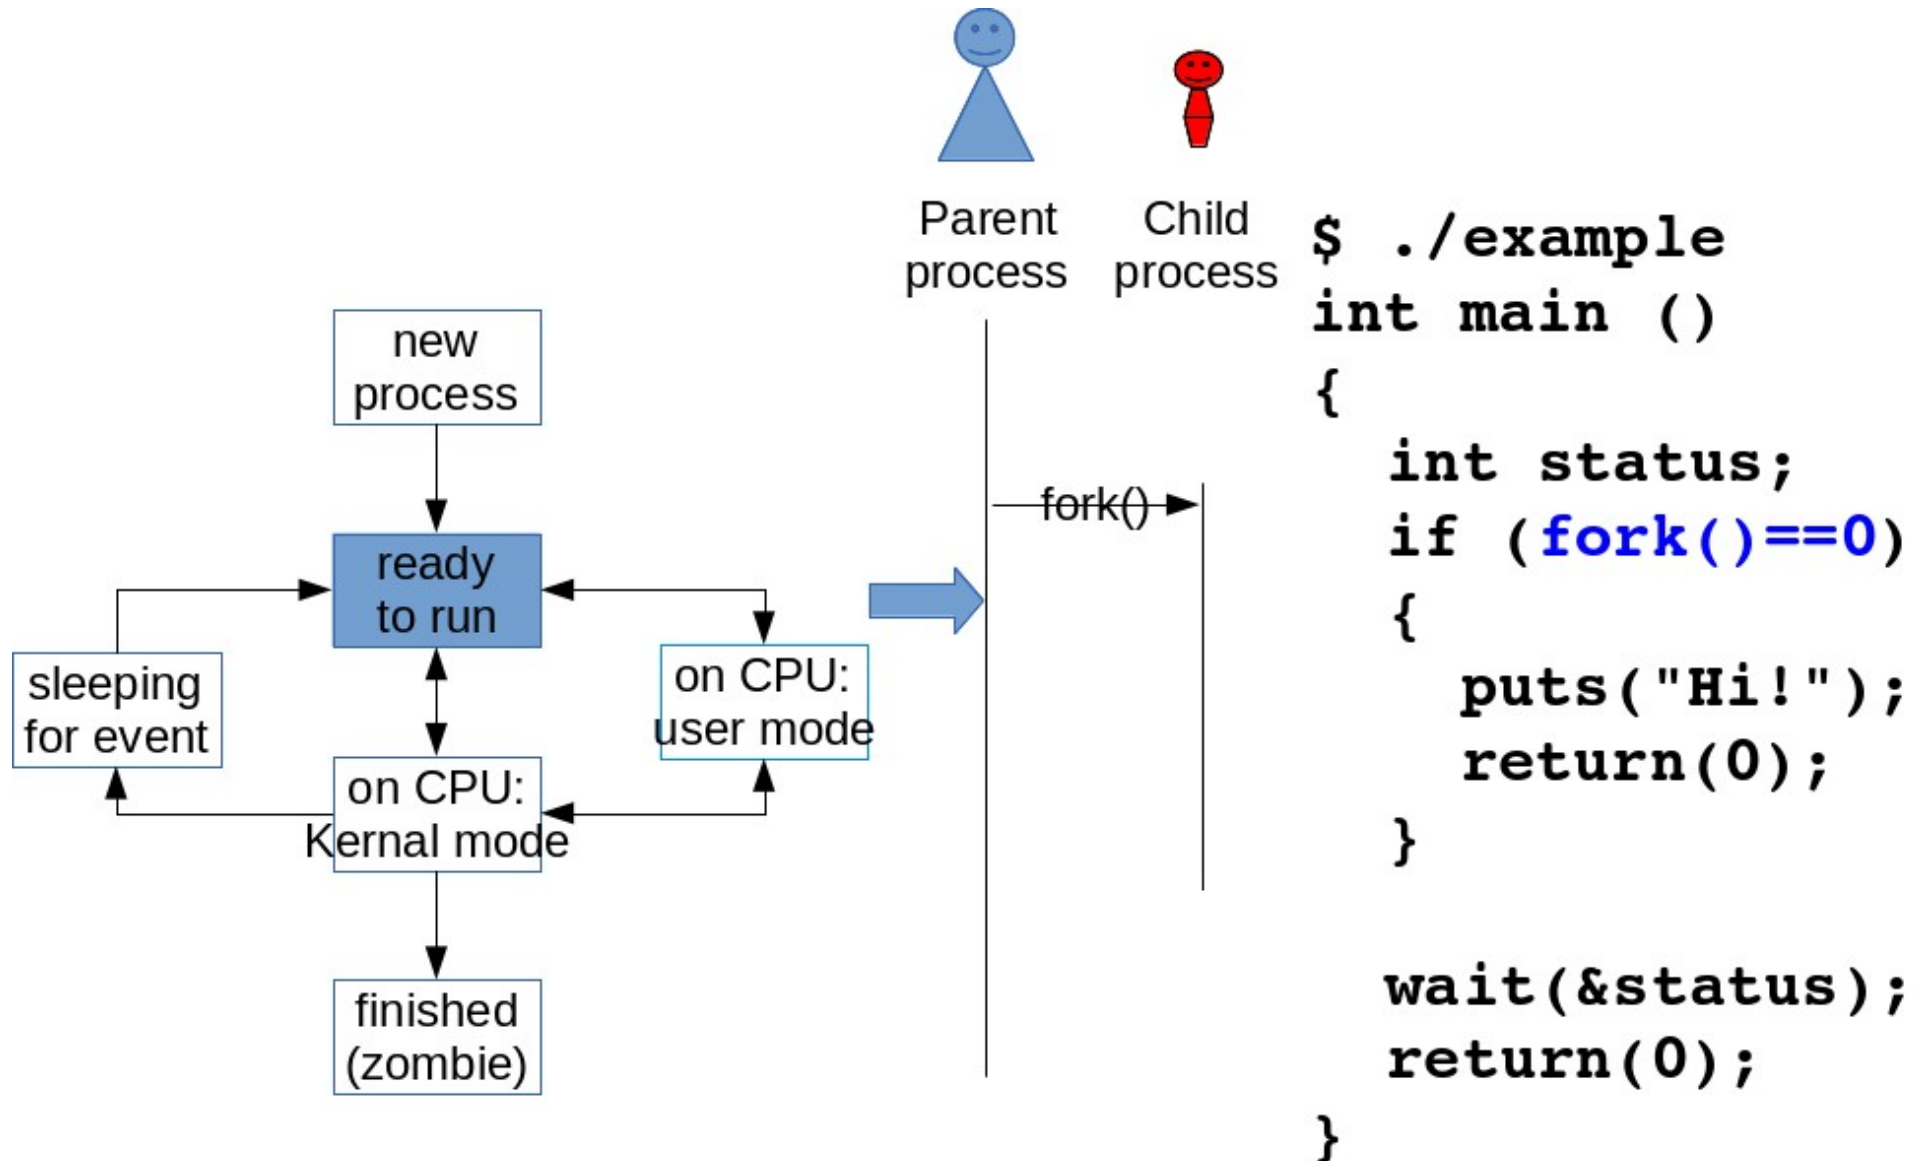
Recall fork system call (2)



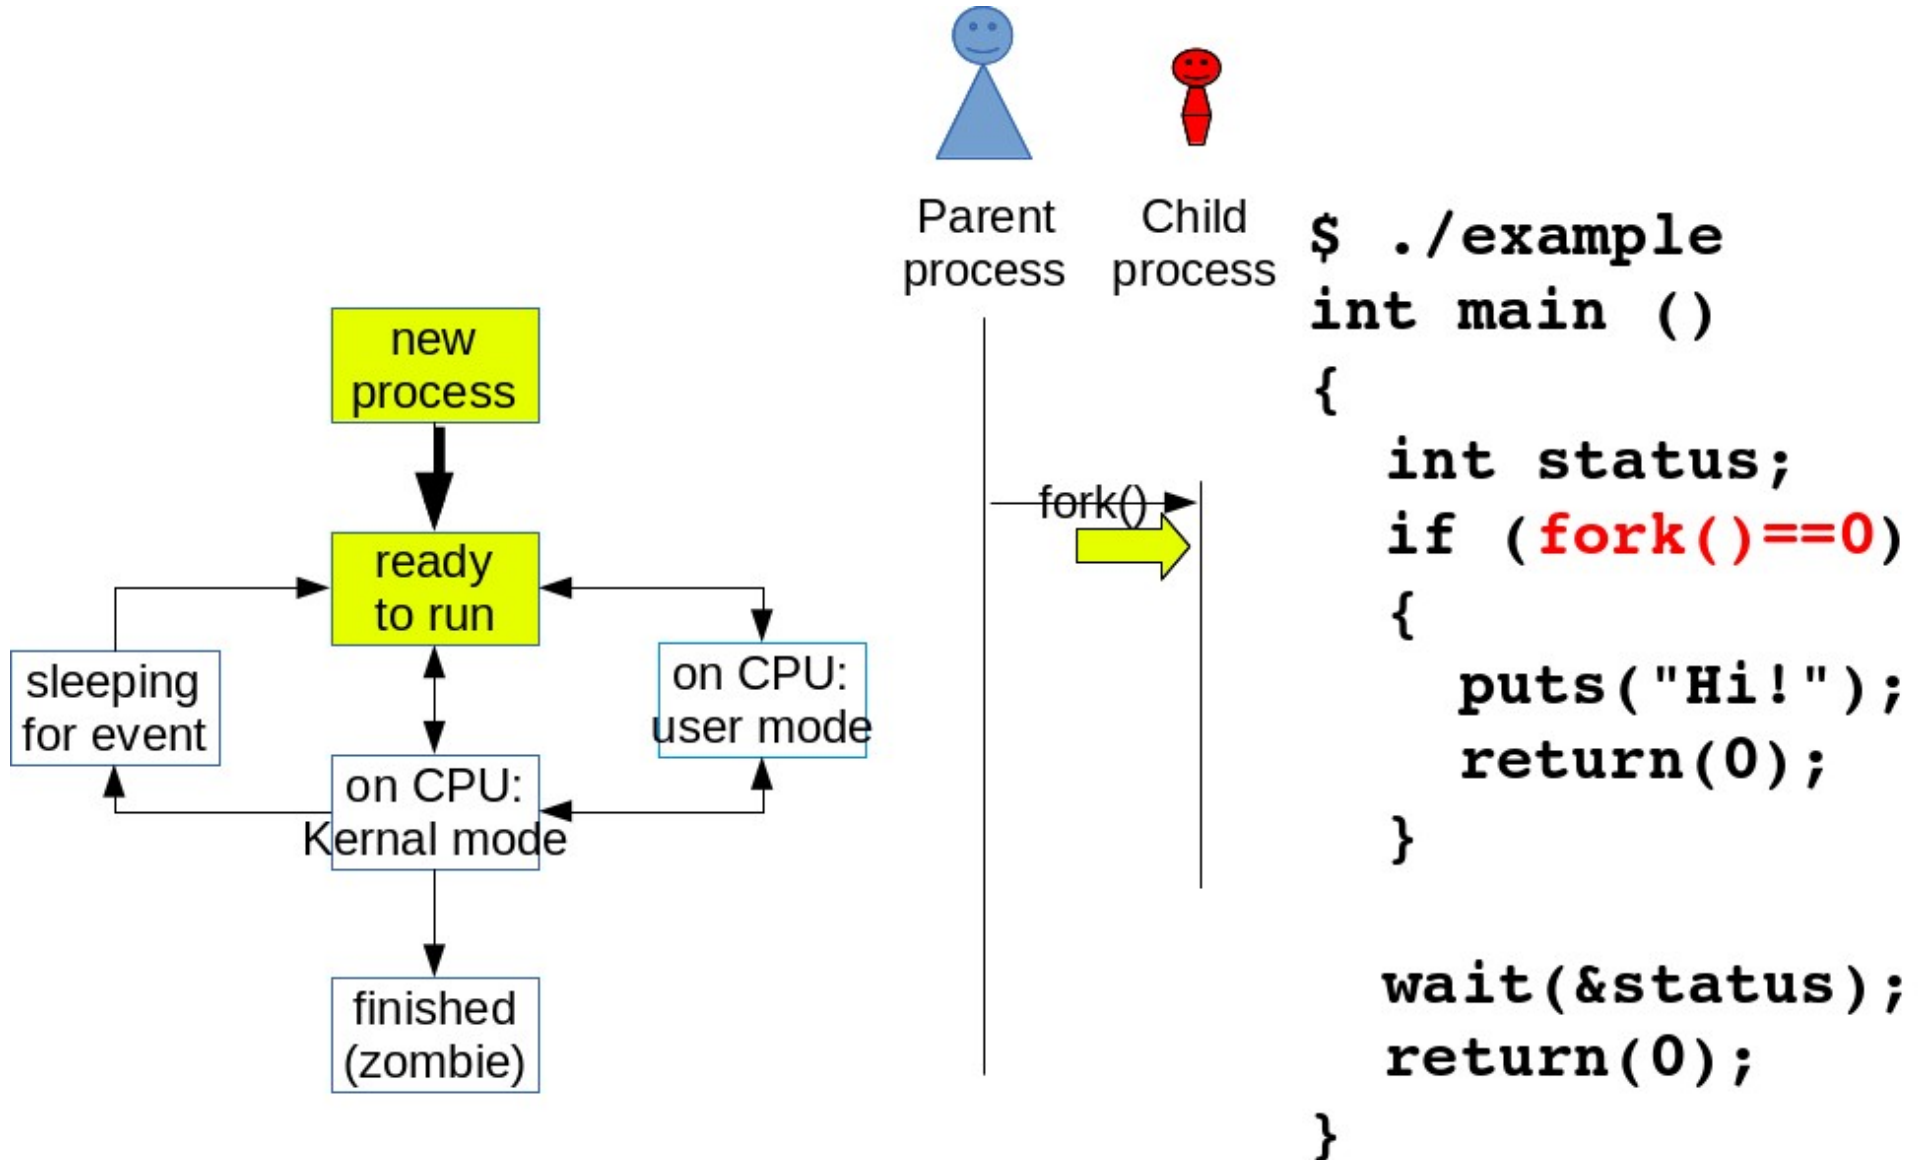
Recall fork system call (3)



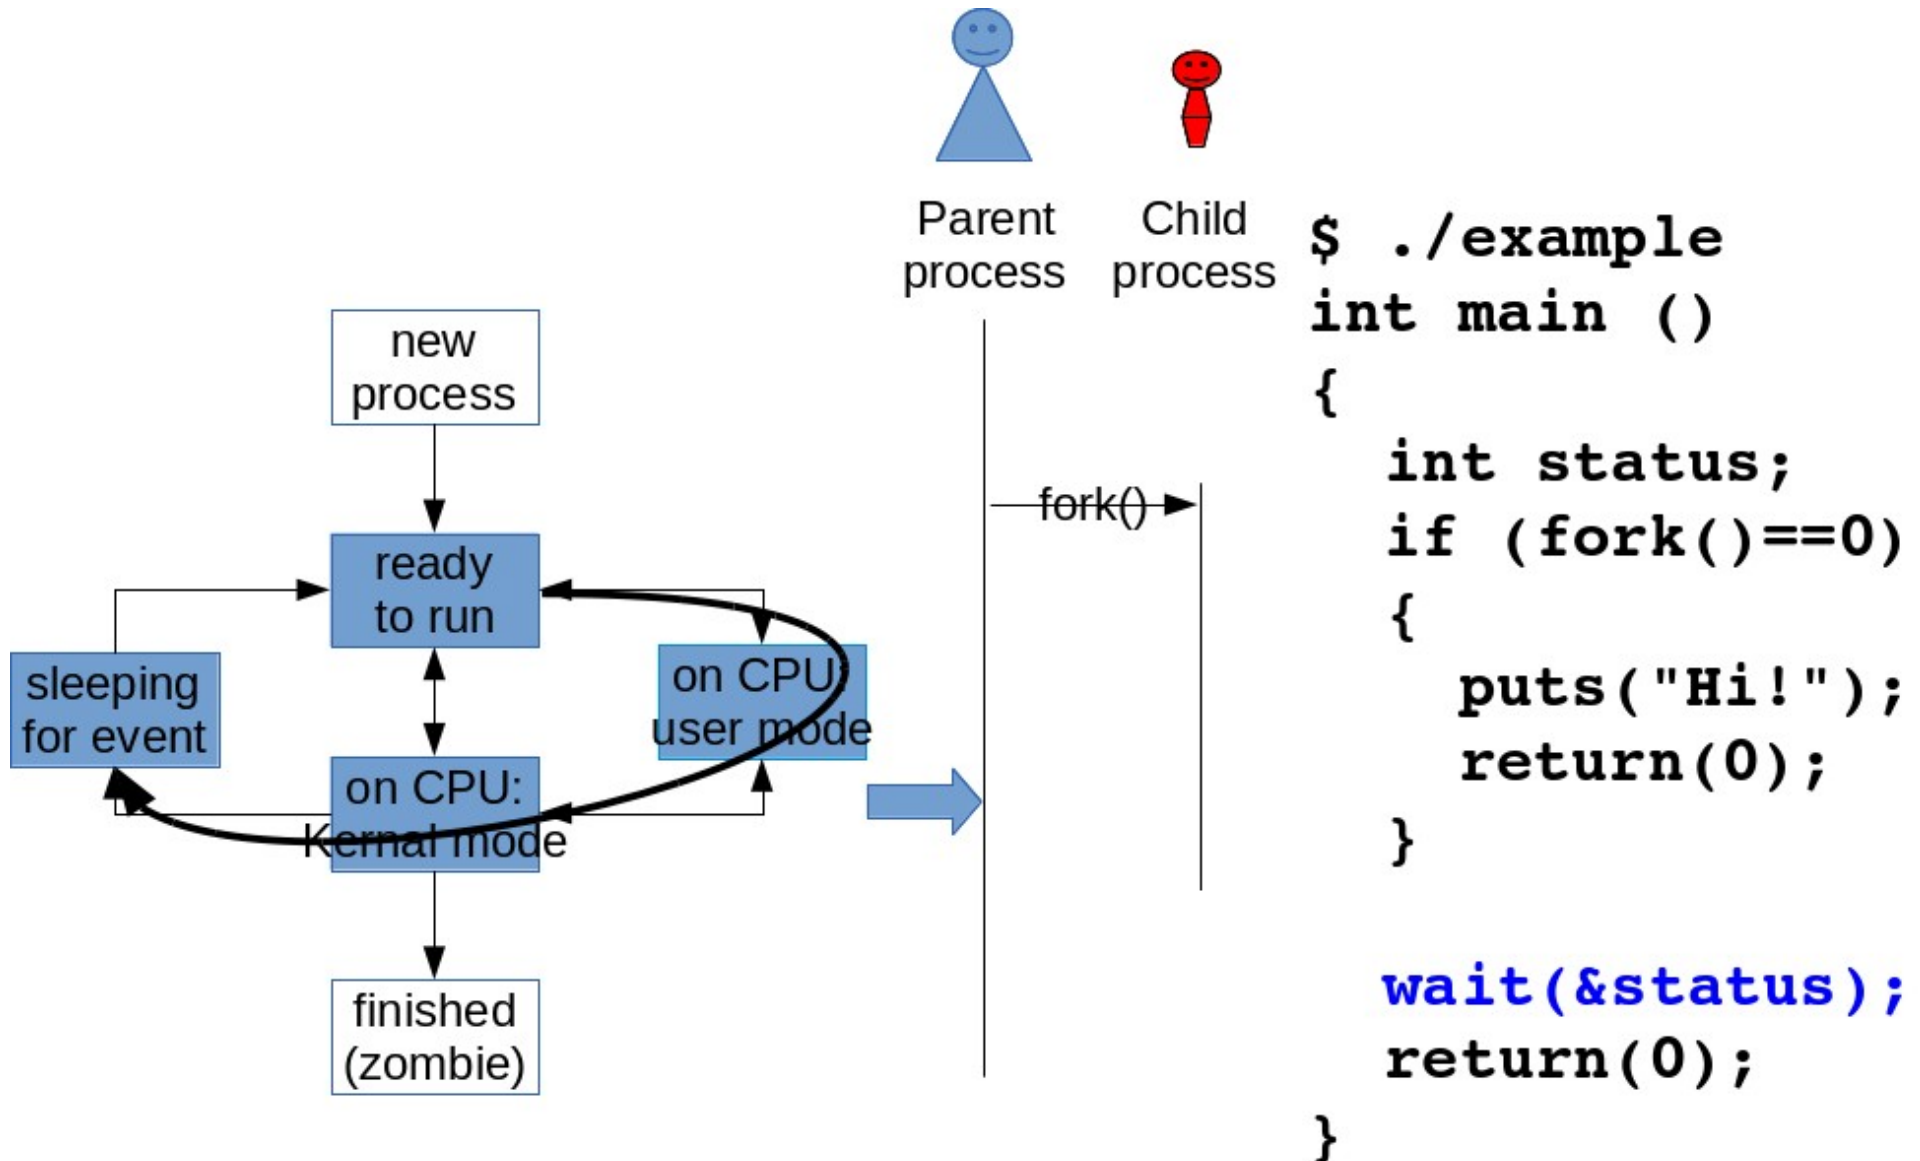
Recall fork system call (4)



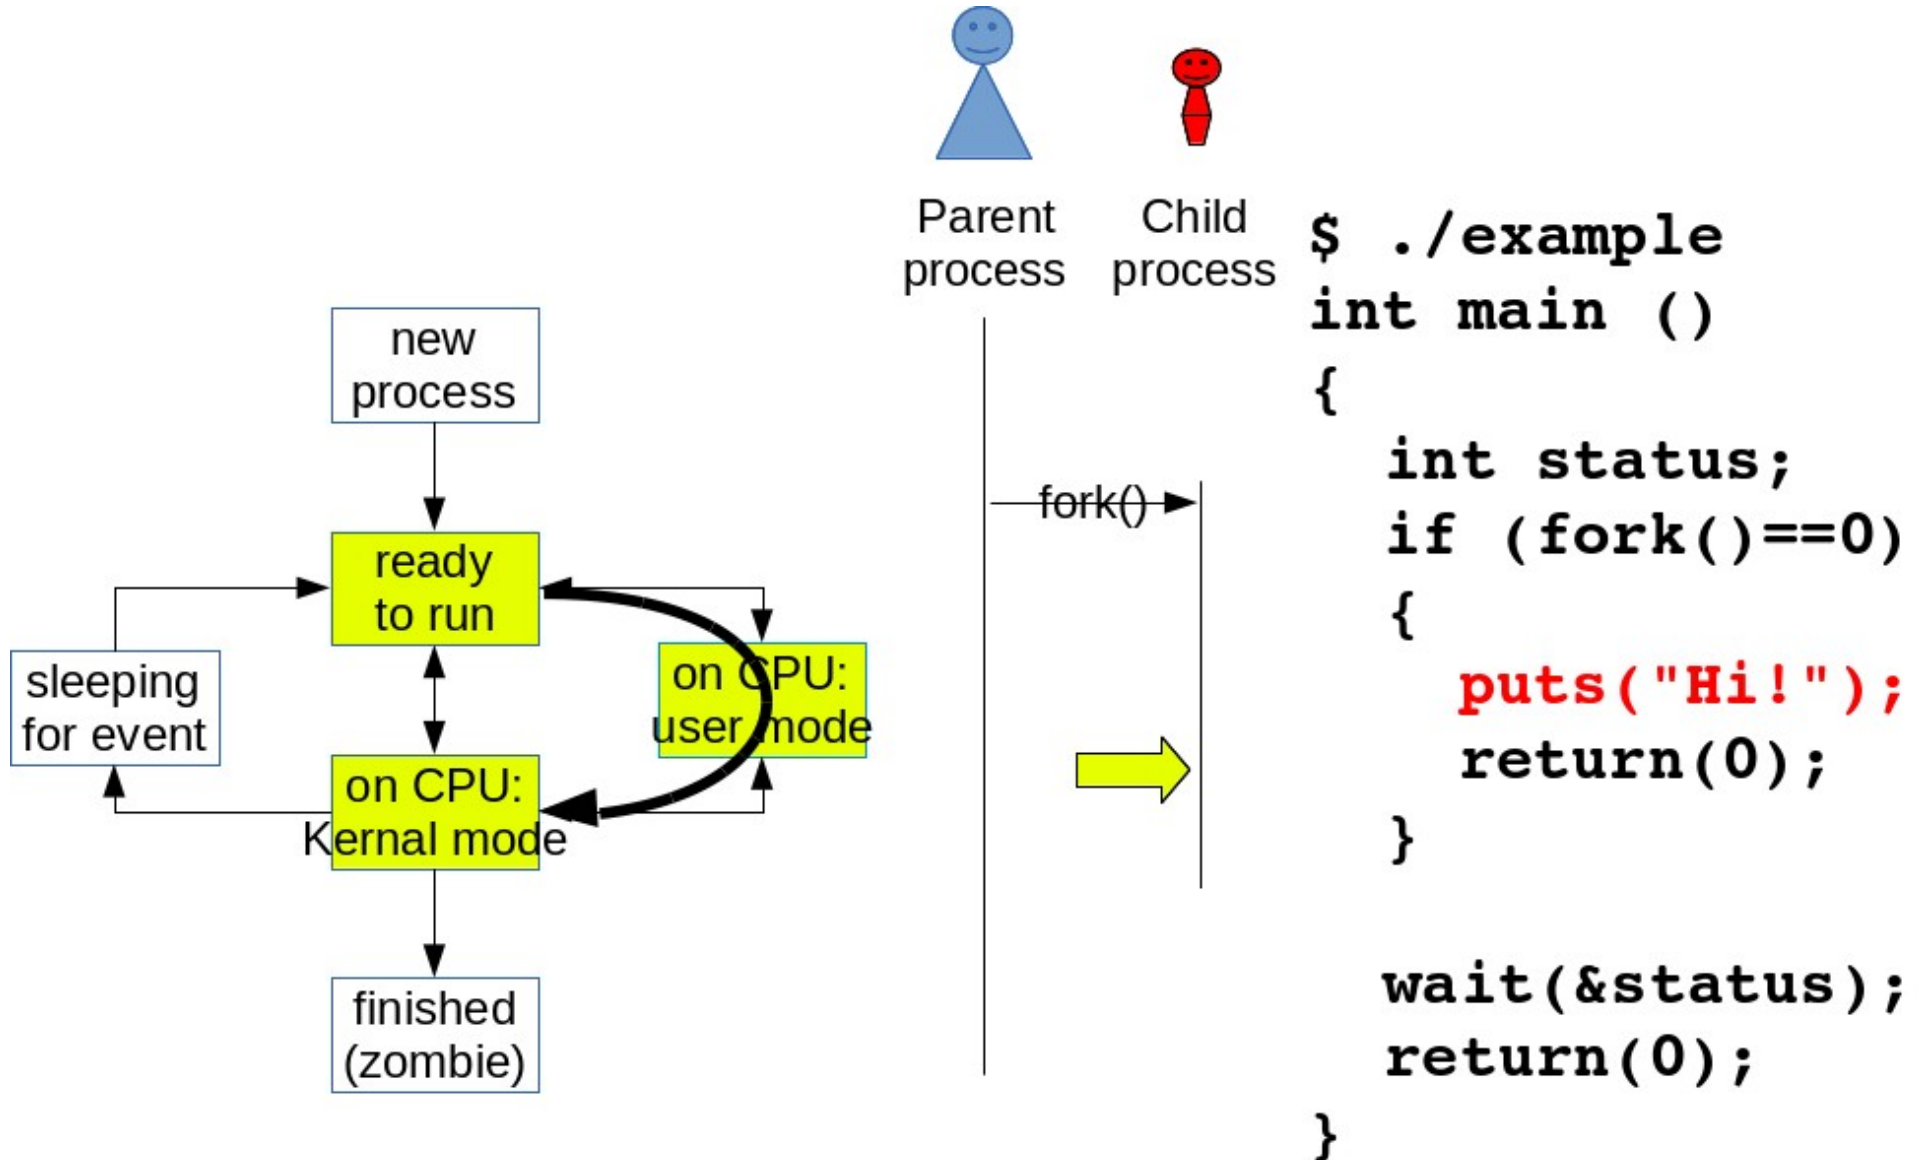
Recall fork system call (5)



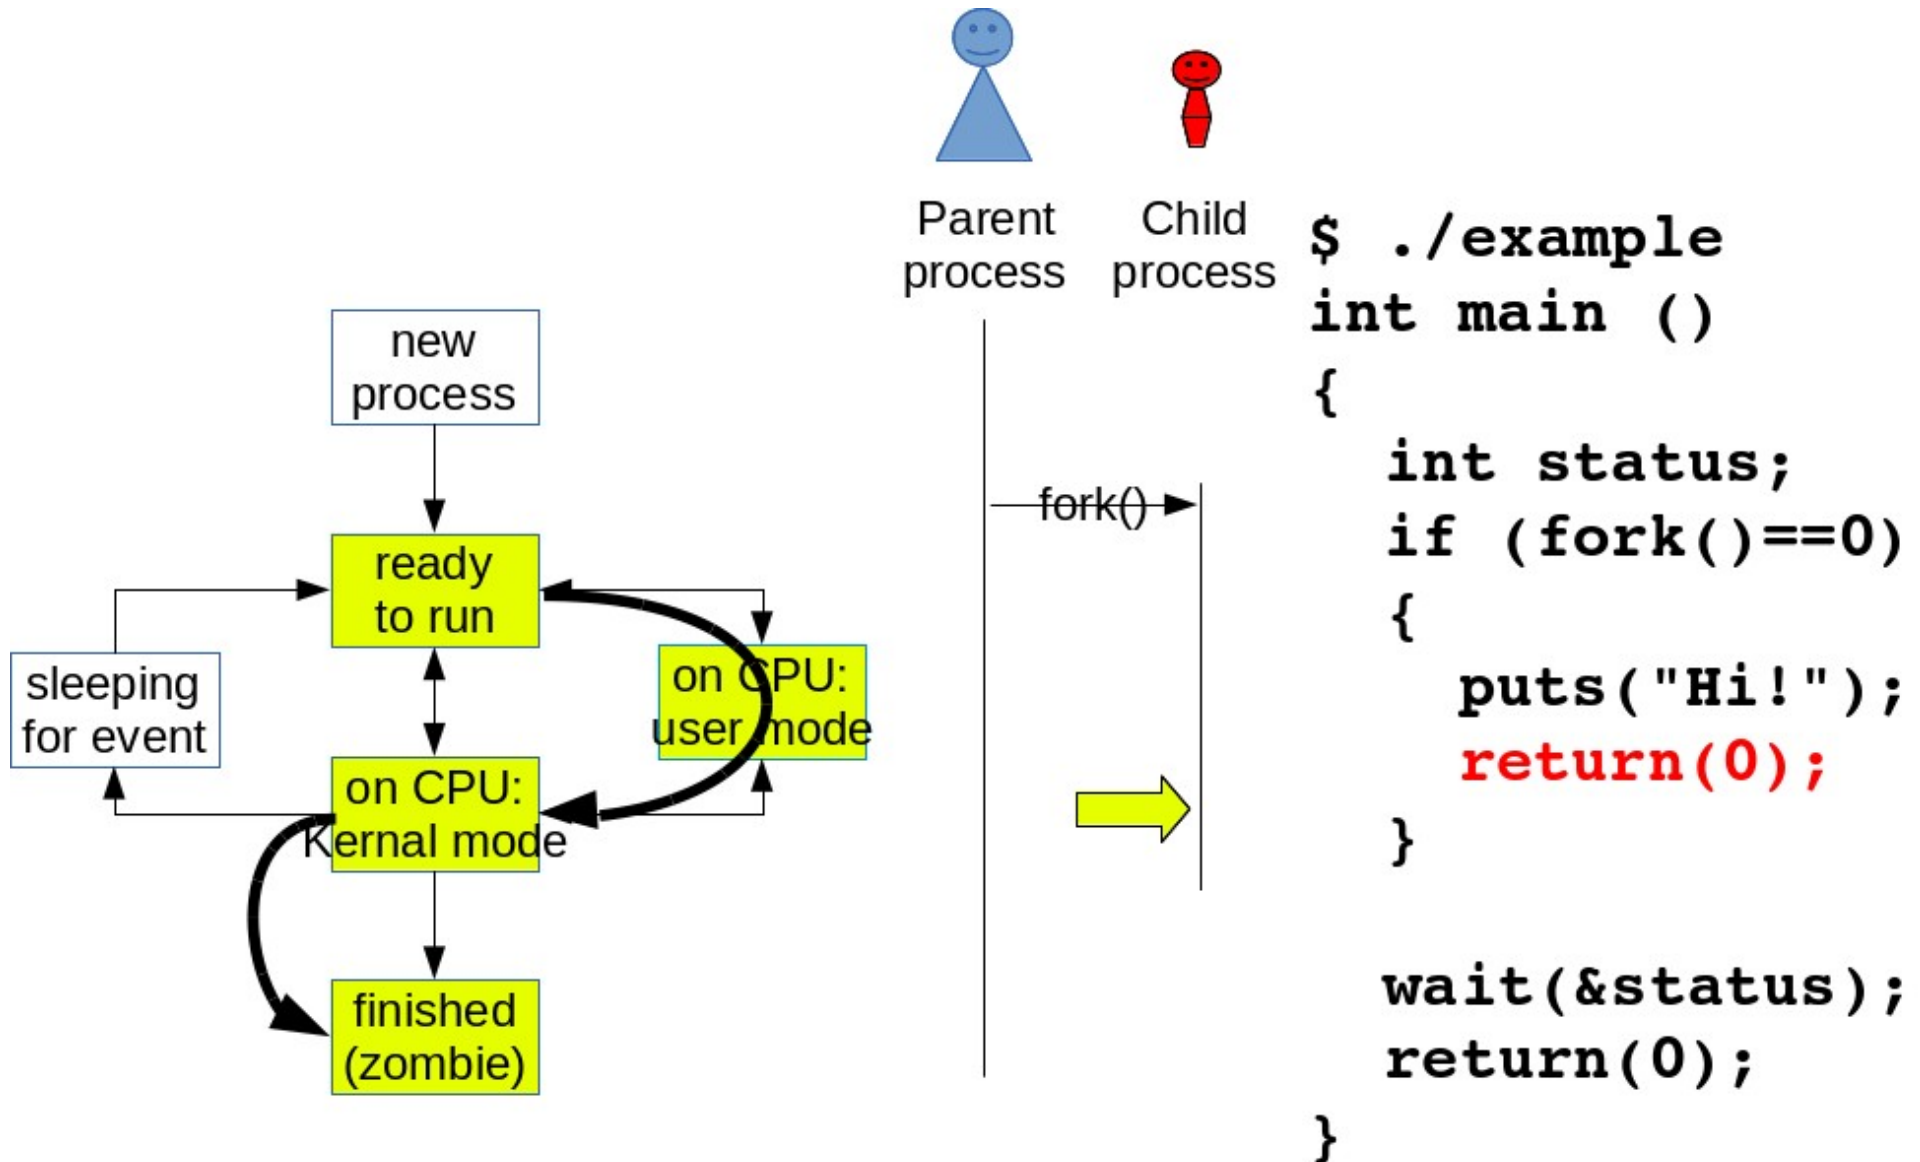
Recall fork system call (6)



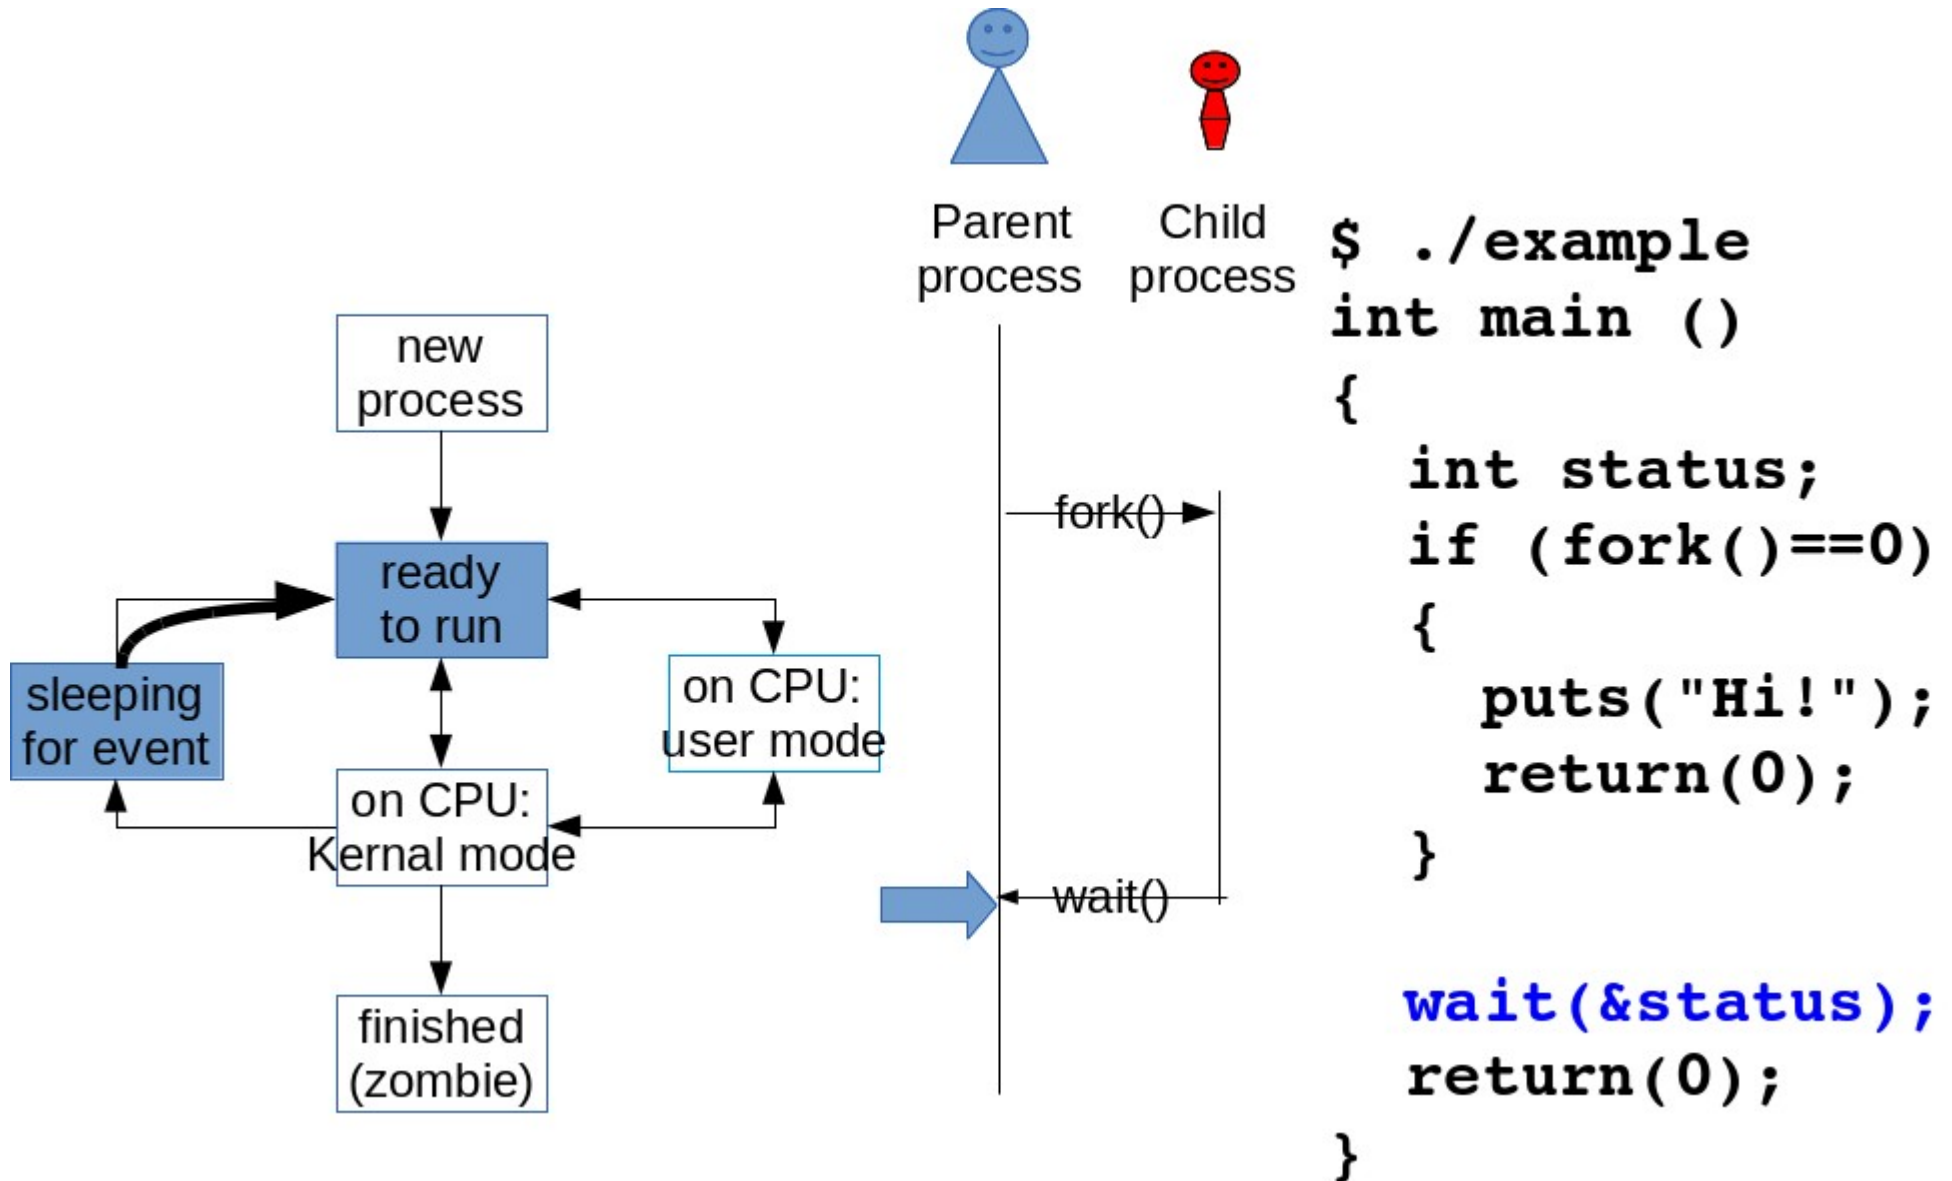
Recall fork system call (7)



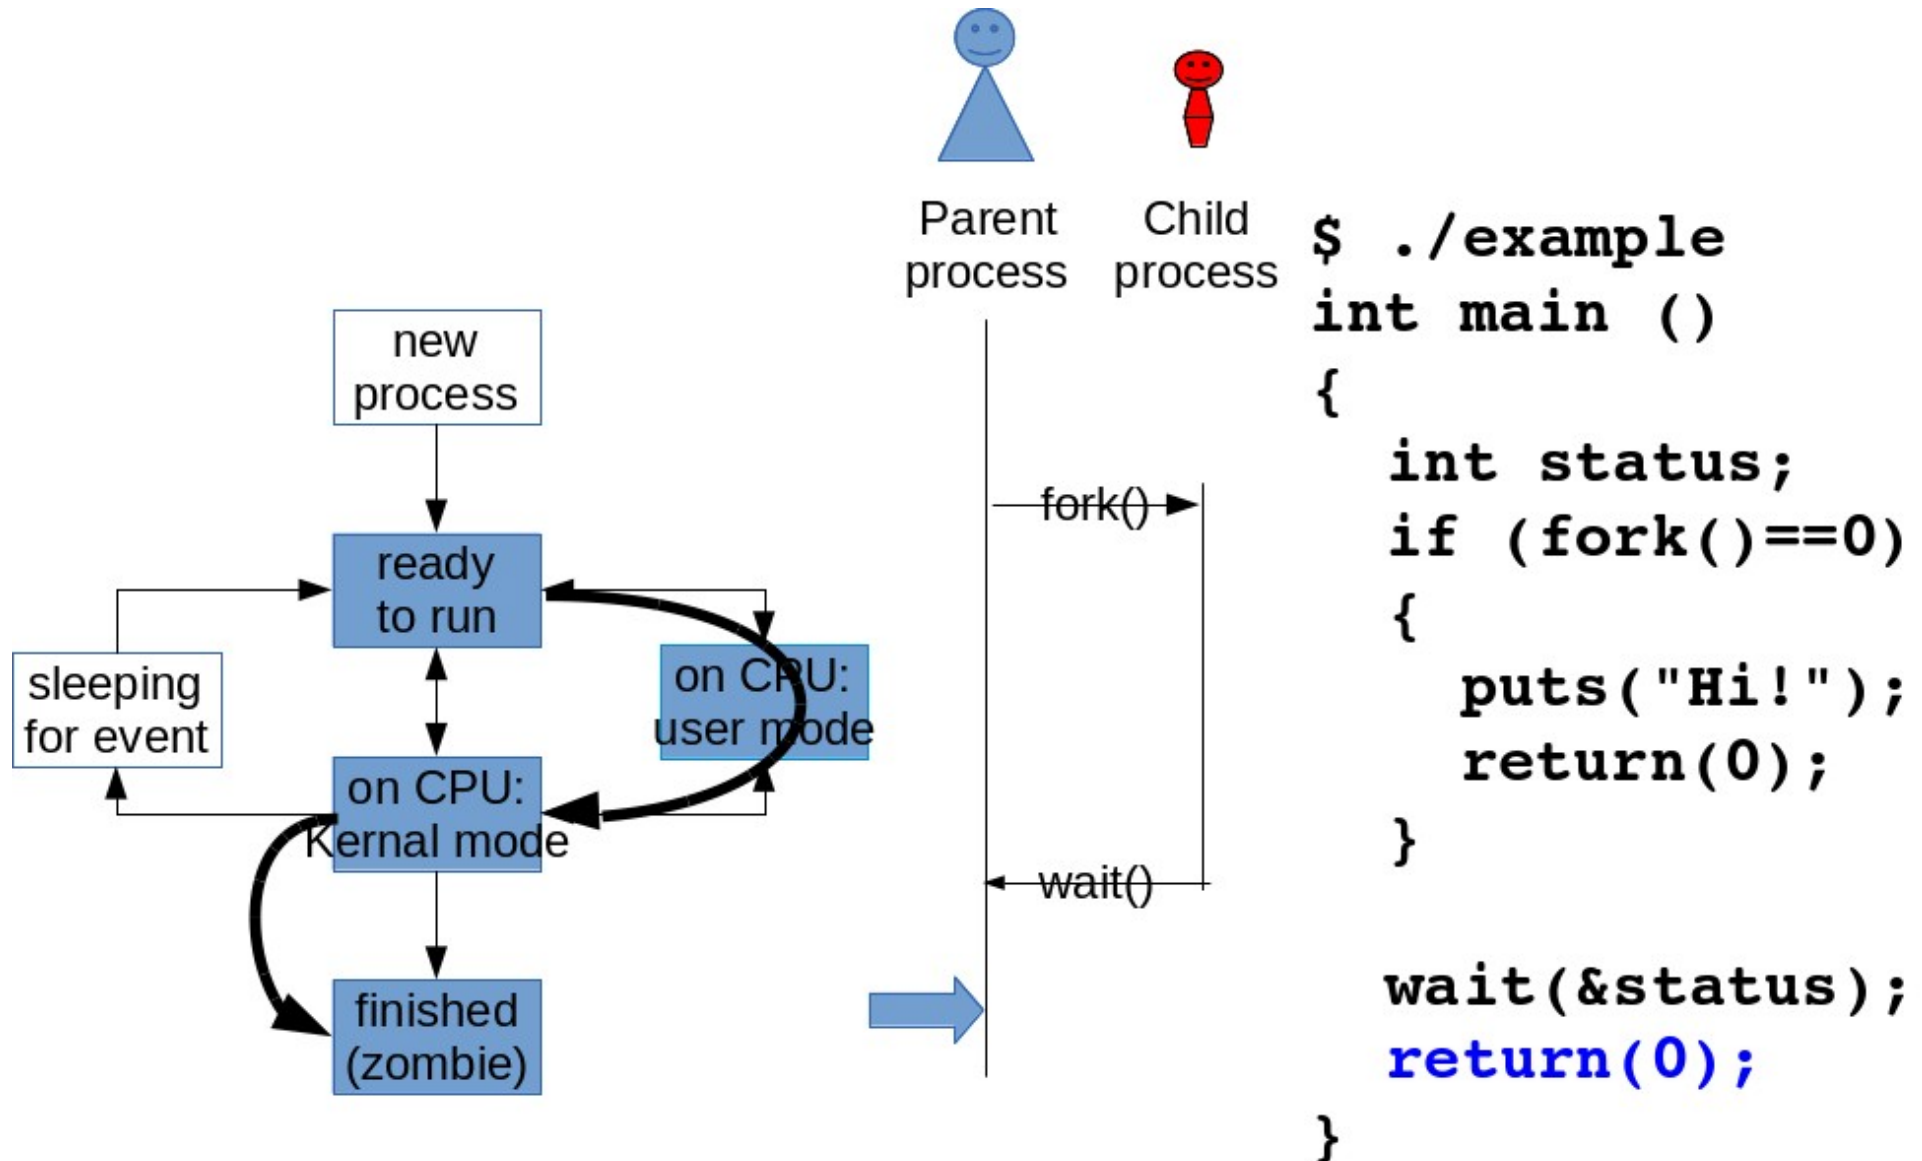
Recall fork system call (8)



Recall fork system call (9)



Recall fork system call (10)




Remember the shell . . .

Let's re-write our simple shell program from last time.

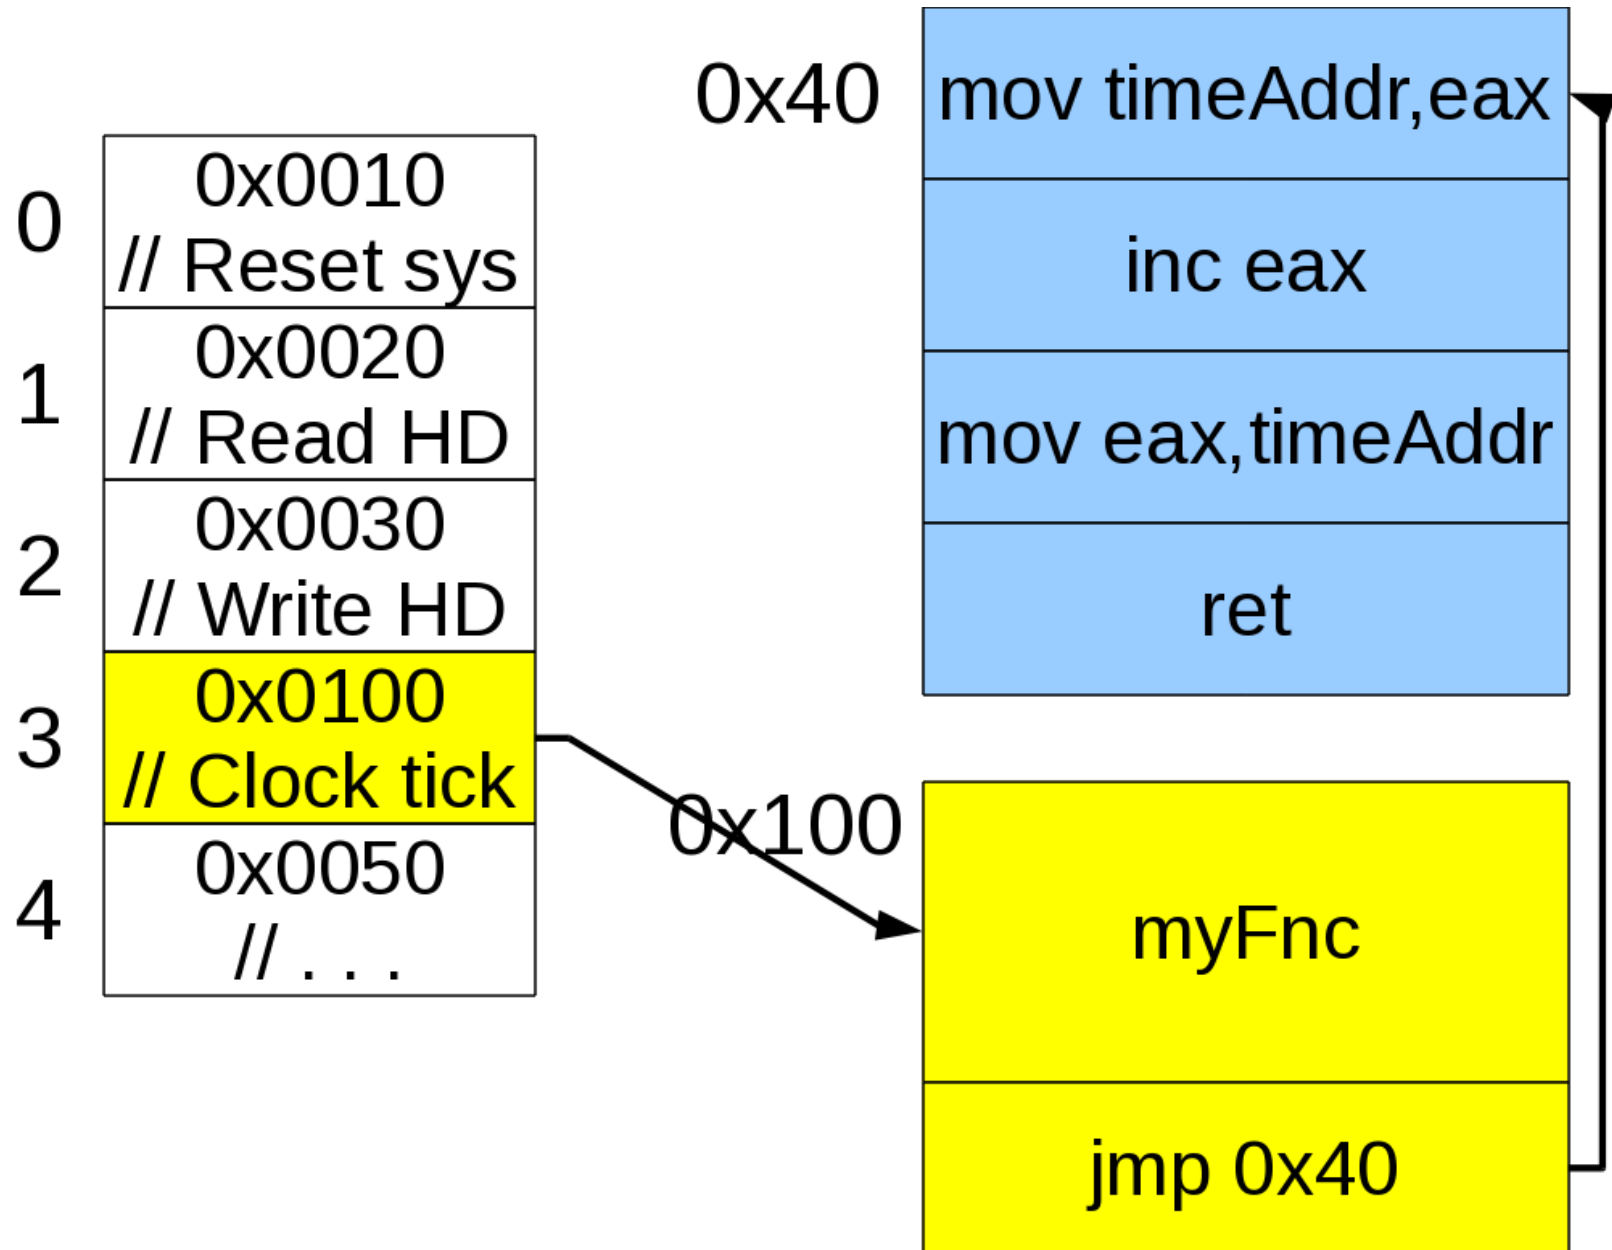
We may use:

- `fork()`
- `execl(char* path, char* arg0, char* arg1, . . .)`
- `wait(int* statusPtr)`
- `waitpid(pid_t pid, int* statusPtr, int flag)`

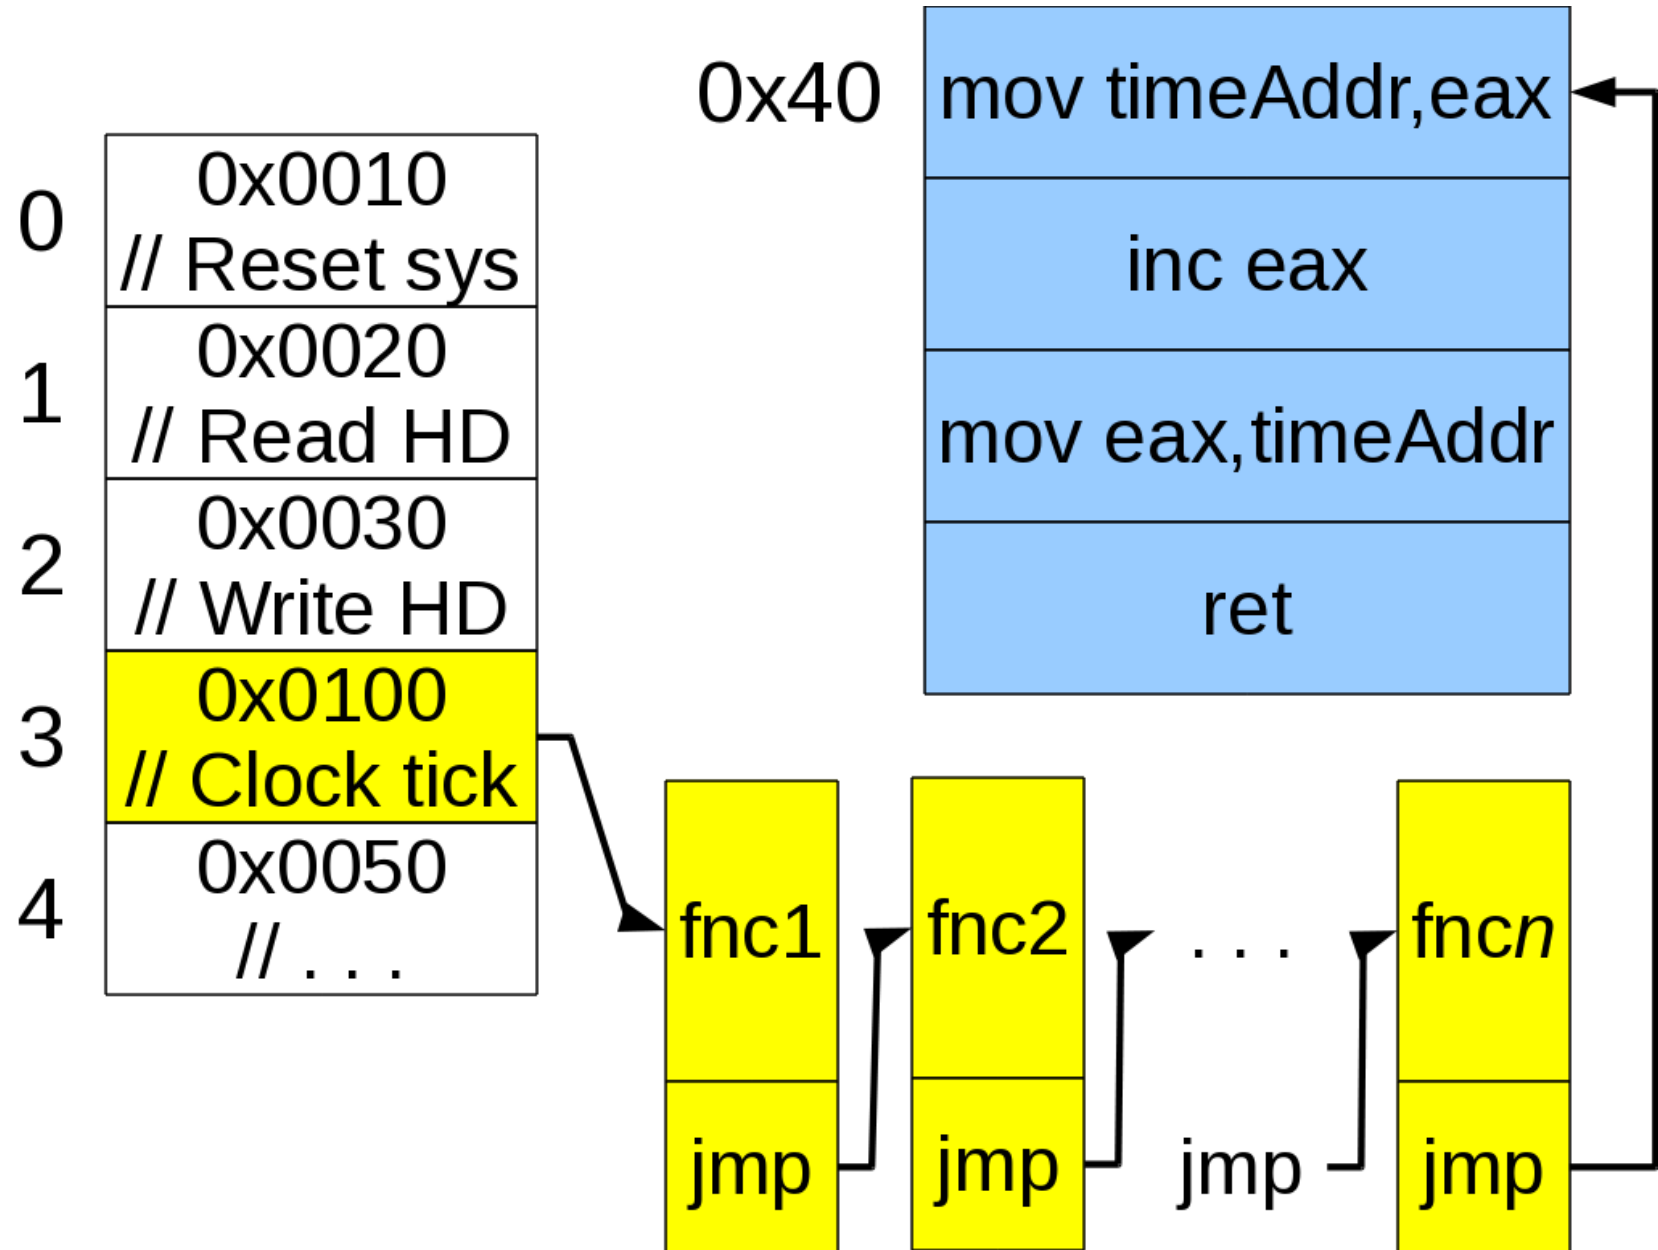
Remember: One interrupt vector table for whole computer



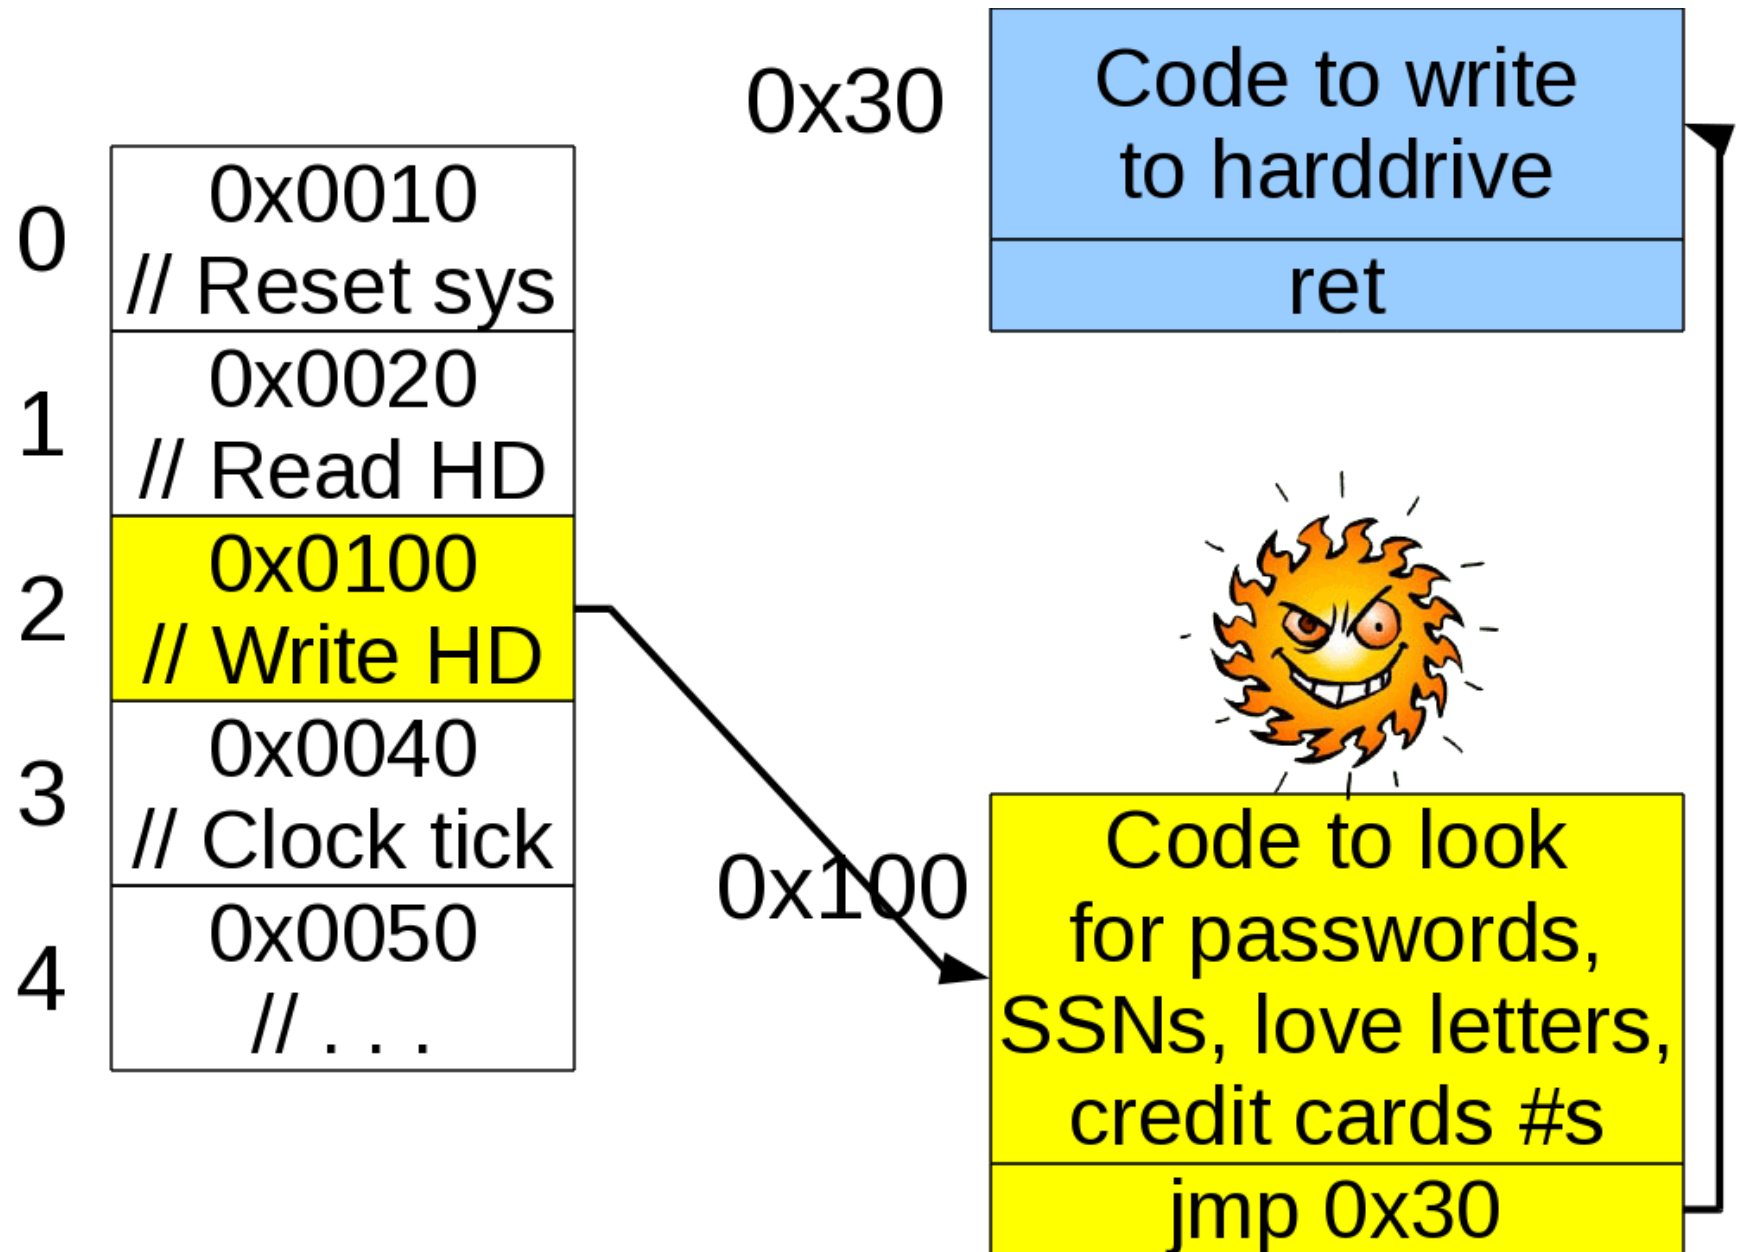
Remember how to do something periodically?



Is this technique efficient when there are many things to schedule?



Is this technique secure?



Signals: a better idea

The whole computer still has its interrupt table

- Nobody can mess with it except OS and ROM BIOS

Each process has its own “personal” interrupt table

It's “interrupts” include

- You set the alarm and went to **sleep()**, now it's *Time to wake up!*
- *Hey!* One of your children just finished!
- Close things down and finish, *or less politely . . .*
- *Die now, Punk!*

List of signals (Linux i386) (1)

The default actions are:

- **Term**: terminate the process
- **Core**: terminate the process and dump the core
- **Ign**: ignore the signal
- **Stop**: stop the process

Signal	Value	Action	Comment

SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process.
SIGINT	2	Term	Interrupt from keyboard (Ctrl-C)
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from abort(3)
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal

List of signals (Linux i386) (2)

Signal	Value	Action	Comment

SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from alarm(2)
SIGTERM	15	Term	Termination signal
SIGUSR1	10	Term	User-defined signal 1
SIGUSR2	12	Term	User-defined signal 2
SIGCHLD	17	Ign	Child stopped or terminated
SIGCONT	18		Continue if stopped
SIGSTOP	19	Stop	Stop process
SIGTSTP	20	Stop	Stop typed at tty
SIGTTIN	21	Stop	tty input for background process
SIGTTOU	22	Stop	tty output for background process


How OS handles signals (1)

Each process has its own signal table:



Some signals may be ignored or handled special

***Signal 1
sent to pid
123***



How OS handles signals (2)

Question: What happens if multiple signals of different number arrive around same time?

Answer: The lowest numbered signal takes precedence

**Signals 0,2
sent to pid
123**



Sorry gotta
Kill you.



How OS handles signals (3)

Question: What happens if multiple signals of same number are pending?

Answer: Because they are not buffered, any called routine will only be called once and should be clever enough to handle any number of waiting cases.

**Signals 2,2,2
sent to pid
123**



Run `fnc()`
once



A look at those actions

Signals so serious that can't be trapped, ignored:

- **SIGSTOP**: Pause! (Restartable with **SIGCONT**)
- **SIGKILL**: **Die now! Punk!**

Other serious signals include:

- **SIGILL**: Illegal instruction (how could this happen?)
- **SIGSEGV**: Illegal memory reference (how could this happen?)

Other signals are less serious

- **SIGCHLD**: A child has finished (Hey, wouldn't **wait()** or **waitpid()** always catch this for us?)
- **SIGINT**: Ctrl-C
- **SIGHUP**: Controlling terminal hung-up (disconnected)

What does all this mean for me?

You the applications programmer get (some) control over your process' personal “interrupt table”

- Do the ***default action*** for this!
- ***Ignore*** that!
- Do it ***my way*** for the other!

sigaction()

(1) Defining a simple action

```
#include <signal.h>
```

```
int sigaction (int signum,  
               const struct sigaction* act,  
               struct sigaction* oldact);
```

Where:

```
struct sigaction  
{  
    void (*sa_handler)(int); // Simple handler  
    void (*sa_sigaction)(int, siginfo_t*, void*);  
                                // Funkier sig handler  
    sigset_t sa_mask; // Sigs to allow in handler  
    int sa_flags; //  
    void (*sa_restorer)(void);  
};
```

Doing it! (1)

```
#include <string.h>    // For memset()
#include <signal.h>

void    simpleHandle (int signalNum)
{
    //    Handling code here
}

int main ()
{
    // Set up struct to specify the new action.
    struct sigaction act;

    memset(&act, '\0', sizeof(act));
    . . .
```

Doing it! (2)

```
// Do this to do the DEFAULT ACTION
act.sa_handler = SIG_DFL; // Handle by default
sigaction(SIGINT, &act, NULL);
```

```
// or, Do this do IGNORE THE SIGNAL
act.sa_handler = SIG_IGN; // Ignore SIGINT
sigaction(SIGINT, &act, NULL);
```

```
// or, do this to HANDLE WITH YOUR FUNCTION
act.sa_handler = simpleHandler;
                // Handle with simpleHandler()
sigaction(SIGINT, &act, NULL);
```

```
...
```

```
}
```

Your turn!

Write a program that

1. Prints “***You can't stop me!
Ngyeah-ngyeah, ngyeah-ngyeah!***”
2. Pauses for 2 seconds
3. Goes back to 1.

Indefinitely, and that ***cannot*** be stopped by **Ctrl-C**.

(***Question:*** Uh-oh! We've created a ***MONSTER!***
How can we stop it?)

Sending signals to processes (1)

Oh no! The program from the previous slide is still running!

- *Ctrl-C* can't stop it!

If we could send **SIGKILL** (9) we could stop it, but how?

```
$ kill -9 <processId>
```

- *Question*: How do we figure out the PID?

Forgive the *BLOODTHIRSTY* name, it should be called “**sendSignal**”

Sending signals to processes (2)

Process groups

- Each process belongs to a process group
- By default it's its parents group
- Can find out group **pid_t getpgrp()** in **unistd.h**.
- Can change group to one's own process id (or that of another process) **int setpgid(pid_t pid, pid_t pgid)**
- Can send signal to all processes in group by making process number negative

```
$ kill -9 -<processGroupId>
```


Back to the actions!

Handlers should have form **void someName (int sigNum)**

- **QUESTION:** Why is the return type `void`?
- **QUESTION:** Why does it take the signal number as the sole parameter?
- **QUESTION:** If it is not OUR CODE that calls the signal handler then how can the signal handler **see** the state of our program? How can the signal handler **change** the state of our program? (Thru **what type** of variable?)

Your turn again!

Same as before, except this time each time the user does press `Ctrl-C` it randomly prints one of the following:

- “***Ouch!***”
- “***Stop that!***”
- “***That hurts!***”
- “***Mercy!***”

HINT: Use `switch(rand() % 4) { . . }` to jump to cases 0 to 3 randomly

Sending signals to processes (3)

Application programs can signal each other too!

```
#include <sys/types.h>
```

```
#include <signal.h>
```

```
int kill (pid_t procToSignal, int sigToSend);
```

Again, forgive the *BLOODTHIRSTY* name, it should be called “**sendSignal()**”

Back to the shell

Our shell program waited for child process to finish.

Revise it so that it may have multiple children (*ie.* one of the children can run in the “background”)

The background child process can finish at any time. (When it does SIGCHLD is sent to the parent.)

Reap the child so it doesn't remain a *ZOMBIE* too long.

HINT: use:

```
action.sa_flags =  
    SA_NOCLDSTOP // No SIGCHLD on stop(pause) or resume  
| SA_RESTART; // If interrupted in sys call then  
                // restart sys call after sig handler
```

Back to the shell (2)

If multiple children finish around the same time
will the different SIGCHLD signals be queued?

How can we revise our previous program to
LEAVE NO ZOMBIES?

- **HINT:** Install this signal handler:

```
void      sigChldHandler      (int      sig)
{ pid_t pid;
  int     s;
  while( (pid=waitpid(-1,&s,WNOHANG)) > 0 )
    if (WIFEXITED(s) != 0)
      printf("%d returned %d\n",
             pid,WEXITSTATUS(s));
    else
      printf("%d crashed!\n",pid);
}
```

sigaction()

(2) Defining more complex actions

```
#include <signal.h>
```

```
int sigaction (int signum,  
               const struct sigaction* act,  
               struct sigaction* oldact);
```

Where:

```
struct sigaction  
{  
    void (*sa_handler)(int); // Simple handler  
    void (*sa_sigaction)(int, siginfo_t*, void*);  
                                // Funkier sig handler  
    sigset_t sa_mask; // Sigs to allow in handler  
    int sa_flags; //  
    void (*sa_restorer)(void);  
};
```

Look at all the goodies in `siginfo_t` (1)

```
siginfo_t
{
    int      si_signo;        // Signal number
    int      si_errno;        // An errno value
    int      si_code;         // Signal code
    int      si_trapno;       // Trap number that caused
                              // hardware-generated signal
                              // (unused on most archs.)

    pid_t    si_pid;          // Sending process ID
    uid_t     si_uid;         // User ID of sending proc
    int       si_status;      // Exit value or signal
    clock_t   si_utime;       // User time consumed
    clock_t   si_stime;       // System time consumed
    sigval_t  si_value;       // Signal value
    int       si_int;         // POSIX.1b signal
    void      *si_ptr;        // POSIX.1b signal
    . . .
}
```

Look at all the goodies in `siginfo_t` (2)

```
    . . .  
int      si_overrun;    // Timer overrun count;  
                        // POSIX.1b timers  
int      si_timerid;    // Timer ID;  
                        // POSIX.1b timers  
void     *si_addr;      // Mem loc that caused fault  
long     si_band;       // Band event (was int in  
                        // glibc 2.3.2 and earlier)  
int      si_fd;         // File descriptor  
short    si_addr_lsb;   // Least sign. bit of addr.  
                        // (since kernel 2.6.32)  
}
```


sigaction() (Example 2a)

```
/* More advanced sig handlers take */
/* tell which process sent the signal */

#include <signal.h>
#include <string.h>
#include <stdio.h>

void    signal_handler
(int sig, siginfo_t* infoPtr, void* dataPtr)
{
    printf("signal:[%d], pid:[%d], uid:[%d]\n",
           sig,
           infoPtr->si_pid,
           infoPtr->si_uid );
    // dataPtr is not used so much
}
```

sigaction() (Example 2b)

```
int main (int argc, char *argv[])
{
    struct sigaction sa;
    memset(&sa, '\0', sizeof(struct sigaction));
    sigemptyset(&sa.sa_mask );

    sa.sa_flags= SA_SIGINFO //Install sa_sigaction
                        // (as opposed to
                        //  sa_handler)
                | SA_RESTART; //If interrupted in
                        // sys call then
                        // restart sys call
                        // after signal handler

    sa.sa_sigaction = signal_handler;
    sigaction(SIGINT, &sa, NULL);
}
```

sigaction () (Example 2c)

```
int i;  
for (i = 0; i < 60; i++)  
{  
    printf("%2d of 60\n",i);  
    sleep (1);  
}  
return(EXIT_SUCCESS);  
}
```

Our turn (1)!

There will be a parent process (“owner”) and a child (“Elmo”).

When Elmo receives **SIGINT**, somebody is trying to tickle it.

Elmo laughs when its owner tickles it (send **SIGINT**), but is weary of tickles from anyone else.

Our turn (2)!

The owner should:

- (1) **fork ()** a child, print its pid and keep it in var.
- (2) Enter a **for (i=0; i<4; i++)** loop where:
 - (a) process waits for user to press enter
 - (b) process sends **SIGINT** to Elmo

Elmo should:

- (1) Install an advanced **SIGINT** handler:
 - (a) If the owner sent **SIGINT**, it prints
“Hee hee, that tickles!”
 - (b) If any other process sent **SIGINT**, it prints
“Elmo does not know you.”
- (2) Does **while(1) sleep(1);** awaiting tickling.

sigaction()

But wait! *There's more!*

Way more detail!

See *\$ man sigaction*

SIGALRM



```
#include <unistd.h>
```

```
unsigned int alarm(unsigned int secs);
```

If **secs**>0 then tells OS “Send **SIGALRM** to me *secs* seconds in the future”

If **secs**==0 then tells OS “Clear any alarms I may have set”

Either case returns number seconds until next (now cleared) alarm, or 0 if there were none.

Your turn!

Write a program that:

- (1) Let's the user type in how many seconds they want to wait until an alarm goes off
- (2) Continually prints “***Tick-tock***” until the alarm goes off. Then it prints “***Ding-ding***” and stops.
- (3) If the user press `Ctrl-C` then it prints how many seconds are left and waits for the user to press `Enter`. Then it goes back to (2).

We don't have time now, but also check out

- **alarm()**: Like **alarm()**, but arguments are in microseconds, not second.
- **sleep()** and **usleep()**: Use **SIGALRM** signal to pause the process for given number of seconds (**sleep()**) or microseconds (**usleep()**)
- **man sigaction**: There is way more details on signals.
- **signal()**: The old-school way to install signal handlers.
- **setjmp()**, **longjmp()**: C's inferior, pre-exception way to recover from errors.

Next time: *Threads!*

sigaction()

(3) Seeing what already is being done

```
#include <signal.h>
```

```
int sigaction (int signum,  
              const struct sigaction* act,  
              struct sigaction* oldact);
```

Where:

```
struct sigaction  
{  
    void (*sa_handler)(int), // Simple handler  
    void (*sa_sigaction)(int, siginfo_t*, void*);  
                                // Funkier sig handler  
    sigset_t sa_mask; // Sigs to allow in handler  
    int sa_flags; //  
    void (*sa_restorer)(void);  
};
```

sigaction () Continued!

Install SIGINT handler if not already ignoring:

```
int    main ()
{
    struct sigaction newAction, oldAction;

    //Define ctrlCHandler that wont block other sigs
    newAction.sa_handler = ctrlCHandler;
    sigemptyset (&newAction.sa_mask);
    newAction.sa_flags = 0;

    // See what is currently done for SIGINT
    sigaction (SIGINT, NULL, &oldAction);

    // Install new handler if not currently ignoring
    if (oldAction.sa_handler != SIG_IGN)
        sigaction (SIGINT, &newAction, NULL);
    . . .
}
```

`setjmp()` and `longjmp()`

Old school C way of doing error recovery

`int setjmp(jmp_buf j)`

- Memorize both position in code (`%eip`) and position on stack (`%esp`) inside buffer `j`.
- First time it's called returns 0 (so you know this is the installation case).

`void longjmp(jmp_buf j, int i)`

- ***Uh-oh!*** We have a hard-to-recover-from error!
- Set `%eax` to `i`, jump back to “safe” state described in `j`.

Use

```
jmp_buf j; // Put in global context
```

```
int main()  
{  
    if (setjmp(j) != 0)  
    {  
        // Handle error  
    }  
    attemptToDoSomethingErrorProne();  
}
```

```
void attemptToDoSomethingErrorProne()  
{  
    if (haveMessedUp == 1)  
        longjmp(j, 1);  
}
```

Abuse

```
/* Question: Will this behave properly? */  
jmp_buf j; // Put in global context
```

```
int main()  
{  
    foo();  
    bar();  
}
```

```
void foo ()  
{  
    if (setjmp(j) != 0)  
    {  
        // Handle error  
    }  
}
```

```
void bar ()  
{  
    if (haveMessedUp==1)  
        longjmp(j,1);  
}
```

Bottom line on `setjmp()` and `longjmp()`

If you use them be sure you `longjmp()` to a function that is still on the stack.

Question: What modern error-trapping technique is in C++, Java, etc. that makes this C construct not necessary?